

Integrating a Modern Knowledge-Based System Architecture with a Legacy VA Database: The ATHENA and EON Projects at Stanford

Aneel Advani, MD,^{1,2} Samson Tu, MS,¹ Martin O'Connor, MSc,¹
Robert Coleman, MS,² Mary K. Goldstein, MD,^{2,3} and Mark Musen, MD, PhD¹

¹Stanford Medical Informatics and ²VA Palo Alto Health Care System

³Stanford University School of Medicine

Stanford University, Stanford, CA 94305-5479

We present a methodology and database mediator tool for integrating modern knowledge-based systems, such as the Stanford EON architecture for automated guideline-based decision-support, with legacy databases, such as the Veterans Health Information Systems & Technology Architecture (VISTA) systems, which are used nation-wide. Specifically, we discuss designs for database integration in ATHENA, a system for hypertension care based on EON, at the VA Palo Alto Health Care System. We describe a new database mediator that affords the EON system both physical and logical data independence from the legacy VA database. We found that to achieve our design goals, the mediator requires two separate mapping levels and must itself involve a knowledge-based component.

INTRODUCTION

One of the major hurdles for clinical application of decision support systems is the ability to use the legacy database infrastructure available in the general clinical application environment. The problem is that frequently there are conflicts between the data models of legacy clinical database systems and the data models assumed by the problem solving methods in the decision-support system.¹ We describe the function of a database mediator that allows a modern component-based medical decision-support system, the Stanford EON architecture,² to use a legacy database, the VA VISTA database.³

The VA VISTA system[†] occupies a unique position as a nation-wide electronic medical record. It offers excellent potential for the deployment of advanced decision support systems in order to effect changes to patient care in real-world clinical settings. However, the legacy databases such as the VISTA in the VA system often combine heterogeneous and inconsistent schemas that have accumulated over time. Moreover, some legacy systems, such as the VA VISTA system use a hierarchical database structure, which presents unique challenges for newer applications that assume

relational data models. The VA VISTA system was developed in the 1970s based on a centralized hierarchical database, called FileMan. FileMan is accessed through the M (MUMPS) language.⁴ The hierarchical model, rather than the current relational model, represented a standard at the time the VISTA system was first developed. This situation presents challenges due to the lack of a high-level query language standard and the more limited expressiveness of the hierarchical database model.

The ATHENA Application

One of the clinical applications that uses the EON decision-support architecture at Stanford is project ATHENA (Assessment and Treatment of Hypertension: Evidence-Based Automation) at VA Palo Alto. The ATHENA project is an evaluation of clinical practice guidelines implemented using automated decision support. The overall task in the intervention is to improve compliance with national hypertension care guidelines.⁵ The system is to be deployed in 11 different clinical sites for greater than 7000 patients who are part of the primary care population of VA Palo Alto. The ATHENA system is one of the first automated knowledge-based DSS applications for guideline-based medical care in the VA system.

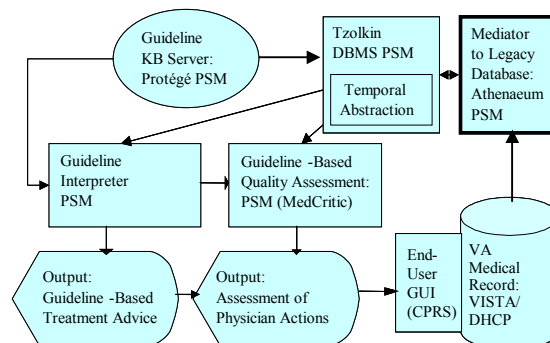


Figure 1 ATHENA system incorporating the EON architecture for component-based decision-support. Each EON component carries out a specific problem-solving task for automated decision-support of guideline-based care.

[†] VISTA also has the older acronym DHCP.

The ATHENA project uses the component-based EON architecture developed at Stanford to automate certain tasks associated with guideline-based care. For each task, there is a specific reasoner or *problem-solving module* (PSM) corresponding to a component in the EON architecture to carry out the task. Such tasks include monitoring the execution of a guideline, using the EON Guideline Interpreter PSM and assessing the quality of the guideline-based treatment, using the MedCritc PSM (see Figure 1).

THE ATHENAEUM MEDIATOR

All the components of the EON architecture use a temporal database management system, called the Tzolkin module, to process queries to an EON-compatible relational database.⁶ However, in the case of the VA hospital environment, the temporal relational data model used in the Tzolkin component must be reconciled with the data model used in the VA medical record database. We have designed a knowledge-based mediator, called Athenaem,[‡] to accomplish this integration. The Athenaem mediator maps one database to another based on the separation of data models into three layers by the ANSI Standards Planning and Requirements Committee (see Table 1).⁷ This methodology allows us to divide the integration of the legacy database into two separate mapping steps (see Figure 2).

In the first mapping, the internal schema of the legacy database is mapped to an application-specific relational schema. This step insulates the application's view of the legacy database from the conflicting physical design of the legacy database, a concept known in the database literature as *physical data independence*.¹ The knowledge-base for this step is a set of declarative frames. These frames describe the semantic transformations from the physical schema (the VA FileMan internal VISTA schema) to a relational schema which includes only those elements needed by the ATHENA hypertension decision-support system. Thus, the external relational schema is physically independent from the internal hierarchical schema.

Second, the external relational schema is mapped to the task-specific conceptual schema, where the task is defined by the functions of the EON PSMs. This transformation insulates the decision-support system from the mismatched semantics of the elements in the legacy database. This concept is referred to as *logical data independence*.¹ The second separate

Table 1 Classification of 3 levels of data models into internal, conceptual, and external schema. Athenaem maps the VA FileMan schema into an intermediate relational schema, and then maps this relational schema to the EON Data Model.

ANSI/SPARC 3-level data model	Database design terminology	Legacy VA database	EON relational database
1. Internal Schema	Physical Design	VA FileMan Database System	Relational Tables
2. External Schema	Specific to Database System Application	ATHENA-Specific Hierarchical Files	EON Temporal Relational Schema
3. Conceptual Schema	Independent of Database but Specific to Task	M Modules (Procedural - not a Data Model)	EON Data Model in Protégé

mapping also allows us to isolate the physical independence from the logical independence of the mapped legacy database. Thus, if we had an alternative method to ensure physical independence from the FileMan architecture, such as using the SQL Interface (SQL-I) system (see discussion below),⁸ we could still re-use the knowledge in the second mapping step.

The Athenaem mediator uses the two mapping steps to asynchronously mediate between Tzolkin and the legacy database. The Athenaem mediator operates on the legacy database at regular intervals, based on the frequency of changes to the relevant patient data. The Tzolkin component then queries the transformed data just as it would for an EON-specific database.

We will use the following example EON Tzolkin query to illustrate the use of the schema and the mapping knowledge-bases in the semantic

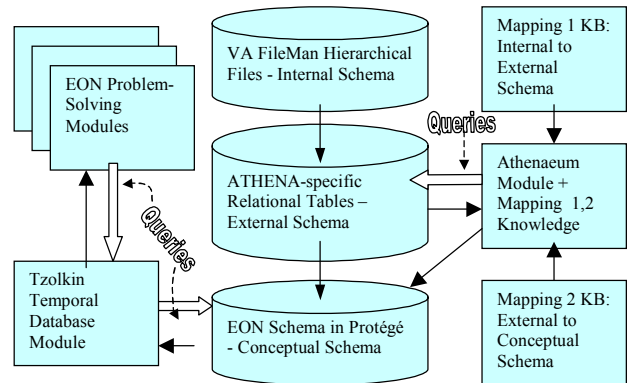


Figure 2 Data Flow Diagram. The Athenaem mediator carries out two separate knowledge-based transformations from: (1) the internal hierarchical schema to external relational schema, (2) the relational schema to the conceptual schema. The block arrows represent query flows, while the straight arrows represent data flows.

[‡] Athenaem is named after the temple of Athena, or "house of wisdom".

transformation from the VA data in order to satisfy the data requirements of the EON PSMs. The query asks when was the last time interval during which patient with id=012345 had primary hypertension that was well controlled for more than six months and during which the patient had a creatinine value of less than 2.1. The query is written in Chronus II, an extended SQL language used in EON to perform temporal abstractions.⁹ It has the form:

```
TEMPORAL SELECT
  LAST problem-name, problem-value
FROM patient-problems-view AS T1
  patient-labs-view AS T2
WHERE patient-id = "012345"
  AND problem-value = "Well_Controlled"
  AND lab-name = "Cr"
  AND lab-value < 2.1
  AND problem-name = "HTN"
WHEN DURATION (T1) > MONTHS (6)
  AND CONTAINS (T1, T2)
```

Note that the query presupposes the EON relational model on which to carry out its temporal select, projection, and temporal join operators.⁹ Since the VA database has a hierarchical internal schema, the Athenaem mediator first maps this to an intermediate relational conceptual schema.

INTERNAL HIERARCHICAL TO EXTERNAL RELATIONAL MAPPING

The main characteristic that distinguishes the hierarchical model from the more modern relational model is its restriction to supporting only one-to-many relations between record-sets, or *files* in the VA VISTA notation.⁷ This model stands in contrast to the relational approach that also supports many-to-one and many-to-many relations between tables directly. Hierarchical databases cannot support these types of relations without computationally expensive work-arounds involving record duplication and other redundancies.

For example, finding all the creatinine measurements over multiple patients is much harder in the hierarchical database. Since the Lab_Data file is a child of the Outpatient Encounter file, which depends on a particular patient encounter, a FileMan search on the Lab_Data file for "all the patients with creatinine < 2.1" would require procedural M coding. The relational model could do this with a simple SQL query on the Current_Labs_and_Vitals table. The relational approach is more expressive and therefore favored by developers. Hence the need for a hierarchical to relational mapping mechanism.

The Athenaem module uses the schema knowledge-base to accomplish this mapping from the physical

design of the hierarchical database to a relational schema. This knowledge-base, built using the Protégé knowledge server tools,² is used to map the relevant hierarchical data files in FileMan to a set of relational tables (see Table 2). For our example query, the information on the creatinine value would have to be obtained from the Lab_Data file and mapped to the Current_Labs_and_Vitals table. The following frame, entered in our schema knowledge-base, allows this mapping:

```
Target: Current_Labs_and_Vitals.Creatinine
Target-type: DECIMAL(4,2)
Operation: CHAR_TO_FLOAT, ALL_ENTRIES
Source: File 63, SubFile 04, SubFile 4
Source-type: VARCHAR(6)
```

The target specifies that the Creatinine field for Current_Labs_and_Vitals hold data in the DECIMAL format from the VA File 63.04.4 that has the creatinine values in the form of characters. The transformations needed include the type change from character to float, and the choice to import all values of creatinine currently in the VA file. Note that this mapping also allows us to import controlled terminology of drugs and diagnoses from the associated files of the VA VISTA database.

Even in this first mapping stage, however, it is clear that several semantic choices have to be made with respect to the relational mapping. For instance, our sample query requires identifying primary hypertension in the patient. However, there were two different potential sources for the current outpatient diagnoses for a given patient: the Purpose_of_Visit file and the Problem_List file. The Purpose_of_Visit file was used for identifying the chief complaint and for billing the visit, whereas the Problem_List file contained a longitudinal problem list that must be kept up to date by primary care physicians. A chart review of 148 paper medical records showed that the Purpose_of_Visit file was 100% sensitive and 79% specific for hypertension, using the paper medical record as a gold standard. The Problem_List was more specific at 95%, but less sensitive at 65%.¹⁰ We made the choice to accept some false positives in order to identify more true positive hypertension patients. Thus, in the ATHENA project, a specific choice was made to point to the Purpose_of_Visit file for identifying patients with hypertension in the schema knowledge base. These semantic mismatches in the "raw" physical files made it necessary to use a configurable knowledge-based approach in creating the relational mapping of the hierarchical database. Hard-coding this mapping would not ensure physical data independence as the national schema recommendations evolved over time.

Table 2 Relational mapping of the hierarchical VA data model. Column on the left indicates the hierarchical files used to create the relational tables, which are listed in the column on the right.

Hierarchical Model VA FileMan Files	Tables in the Relational mapping of VA Files
•Hospital_Location •Outpatient_Encounter	Current_Encounter – date, location, provider
•Purpose_of_Visit •Outpatient_Diagnosis Outpatient_Prescription	Active_Problem_List –noted at time of visit Current_Prescriptions – both active and allowed
•Lab_Data •Vitals_Measurement	Current_Labs_and_Vitals – inpatient and outpatient
•Patient File	Pt_Demographics
•Drug_File	Drug Names Vocabulary
•ICD_Diagnosis	Diagnosis Codes/Names
•Labs Data Dictionary	Lab Tests Field Codes

EXTERNAL RELATIONAL SCHEMA TO EON CONCEPTUAL SCHEMA MAPPING

In the first mapping, we were successful in removing the physical data dependence in the semantics of the VA legacy database. However, there are additional semantic mismatches that exist at the conceptual or logical level of the ANSI hierarchy presented in Table 1.¹ These include mismatches in scope restriction, domain semantics, temporal basis, temporal granularity, and temporal abstraction (illustrated in Table 1). Therefore, the Athenaeum mediator includes a mapping knowledge base to create another set of relational tables. These tables, in addition to being internally consistent, are also semantically consistent with the EON data schema. Table 3 shows how each of these semantic mismatches occurs in the example query presented above.

For instance, the query looks for records where “lab-name ‘Cr’ AND lab-value < 2.1”. The VA Lab_Data file, and hence also the VA Current_Labs_and_Vitals table, include both inpatient and outpatient laboratory results. However, the EON Guideline Interpreter does not filter out values that are inpatient measurements. Knowledge to accomplish this filtering is not embedded anywhere in the database itself. This *scope restriction* mismatch is resolved by encoding the following frame in the mapping knowledge-base:

Target: Patient_Labs_View.Creatinine
Operation: SCOPE_RESTRICTION(Filter_Outpatient)
Source: Current_Labs_and_Vitals.Creatinine

where Filter_Outpatient is a scope restriction function. The Athenaeum module, using a Chronus II query to the Tzolkin module, evaluates it:

SELECT Current_Labs_and_Vitals.Creatinine

FROM Current_Labs_and_Vitals AS T1,
Outpatient_Encounter AS T2
WHEN CONTAINS (T2, T1)

This scope restriction must be used so that we do not use inpatient laboratory data for reasoning about guidelines that only consider outpatient diagnoses. The interpretation of a creatinine value > 2.1 might change in the context of an acute hospital illness.

Similarly, the first condition in the WHERE clause, “problem-name = ‘HTN’” refers to a problem called “HTN” (for hypertension) and not to an ICD-9 code. The diagnostic information from the Purpose_of_Visit file in the VA database comes in the form of ICD-9 codes. Thus, the use of “HTN” as a well-defined problem name depends on a knowledge base that transforms the domain semantics of the “raw” diagnostic codes into well-formed elements of a patient problem list. In this case, our knowledge base used the ICD-9 codes 401.1 and 401.9 for primary and unspecified hypertension respectively as comprising the patient problem of HTN. This excluded the codes 401.0 for malignant HTN and 402.x for HTN with concomitant heart disease from the category. The particular domain semantic transform chosen for the mapping knowledge-base was dictated by the inclusion criteria set out by the national standard guideline for the management of primary HTN.

The reasoning steps described above require a separate application-specific mapping knowledge-base and a mediator that can use the mapping knowledge. We separate the knowledge required for transformations to a relational schema from the application-specific mapping to the EON schema. Procedurally, we use the Chronus II query language on the intermediate relational schema to carry out this

Table 3 Examples of semantic transformations required from VA to EON schemas that are required to answer the example query. Note: HTN refers to hypertension.

EON Module	EON Data Query	VA Data Model	Semantic Transform
Guideline Interpreter	Creatinine - Outpatient Labs	Creatinine – Inpatients & Outpatients	Scope Restriction
Guideline Interpreter	Hypertension Problem Name	ICD-9 Code 401.1	Domain Semantic Transform
MedCritic	Creatinine - Transaction Time	Creatinine - Valid Time	Temporal Basis
Tzolkin DBMS	Time Grain of Month	Timestamp Day/Time	Temporal Granularity
Tzolkin DBMS	HTN = “Well Controlled”	BP values over time	Temporal Abstraction

second mapping. We can thus carry out more complex temporal semantic transformations than would be possible using only standard SQL queries.

DISCUSSION

The functions supported by the Athenaeum mediator include schema integration, semantic mappings, and vocabulary/data dictionary (meta-data) integration. However, unlike some other database mediator architectures,¹¹ that use a generic relational model of medical care, our approach uses an application-specific knowledge-base optimized to each guideline task performed by components in the EON architecture. We also use the Chronus II query language to find the results of semantic transformations that must be applied as a result of the mapping knowledge.

This approach allows us to avoid the pitfalls of overly general approaches to database mediation. A case in point is the standard relational mapping of the VISTA system, the SQL Interface (SQL-I) system.¹² As we saw in the discussion above, the more than 400 record-sets or files that make up the VISTA system frequently contain inconsistencies in semantic scope and conflicting assumptions in domain semantics. To use the SQL-I system as the only relational model for the VA databases would be quite inadequate. Moreover, to actually operate on the VA databases, one has to use an additional proprietary M-to-SQL product. These mappings have their own proprietary semantic assumptions about the most appropriate mapping from the VA database to a relational model. Our knowledge-based approach allows the application-specific knowledge to be separated from the question of relational mapping. Moreover, the schema knowledge-base also removes the internal semantic inconsistencies that the SQL-I relational mapping simply propagates. This allows us to achieve true physical data independence that cannot be achieved using the combination of SQL-I and a proprietary M-to-SQL product.

Thus, a straight generalized relational mapping, without a properly constructed semantically consistent knowledge base cannot be used for automated decision-support. Our approach allows a knowledge-based mediator to represent and automatically resolve inconsistencies and semantic domain constraints that can only be detected by human experts familiar with the guideline application and the assumptions of the VISTA system. Secondly, since these knowledge-bases are configurable, when the underlying VA database

changes in some way, only the knowledge-bases and not the mediator code itself need to be updated.

Acknowledgements

This work was supported in part by NLM grants LM05708 and LM06245 and VA grant HSR&D CPG-97-006. We thank Ravi Shankar, and Larry Lou for programming assistance on this project.

References

- ¹ Elmagarmid A, Rusinkiewicz M, Sheth A, eds. Management of Heterogeneous and Autonomous Database Systems. San Francisco: Morgan Kaufmann; 1999.
- ² Musen MA, Tu SW, Das AK, Shahar Y. EON: A Component-Based Approach to Automation of Protocol-Directed Therapy, JAMIA 1996;3(6):367–88.
- ³ Beattie, MC, Buxton, EC, Wiederhold, V. VA Databases Resource Guide, Vol. V, DHCP V3.0. Palo Alto: VAPAHCS, HSR&D Center for Health Care Evaluation; 1996.
- ⁴ VISTA Software Development Team, Dept. of Veterans Affairs. VA FileMan v. 21.0 Programmer Manual; 1997. URL: www.vista.med.va.gov/softserv.
- ⁵ National High Blood Pressure Education Program. The Sixth Report of the Joint National Committee on Detection, Evaluation, and Treatment of High Blood Pressure. Washington: NIH; 1998.
- ⁶ Nguyen JH, Shahar Y, Tu SW, Das AK, Musen MA. A Temporal Database Mediator For Protocol-Based Decision Support. 1997 AMIA Annual Fall Symposium, Nashville, TN, 298-302. 1997.
- ⁷ ANSI Study Group on Data Base Management Systems, Interim Report. Bulletin of ACM SIGMOD (FDT); 1975; 7(2).
- ⁸ VISTA Software Development, Dept. of Veterans Affairs. VA FileMan SQL Interface (SQLI) Site Manual; 1997. URL: www.vista.med.va.gov/softserv.
- ⁹ O'Connor MJ, Tu SW, Musen MA. Applying Temporal Joins to Clinical Databases. 1999 AMIA Annual Fall Symposium, paper in this volume.
- ¹⁰ Szeto H, Goldstein MK. Accuracy of Computer-Identified Diagnoses in a VA Medical Clinic. Presented at the 17th Annual VA HSR&D Service Meeting, Washington D.C. Feb 24-26, 1999
- ¹¹ Sujansky W, Altman RB. An Evaluation of the TransFER Model for Sharing Clinical Decision-Support Applications. 1996 AMIA Annual Fall Symposium, Washington, D.C.; 1996; 468-472.
- ¹² VISTA Software Development, Dept. of Veterans Affairs. VA FileMan SQL Interface (SQLI) Site Manual; 1997. URL: www.vista.med.va.gov/softserv.